

Assignment Code: DA-AG-010

Regression & Its Evaluation | Assignment

Instructions: Carefully read each question. Use Google Docs, Microsoft Word, or a similar tool to create a document where you type out each question along with its answer. Save the document as a PDF, and then upload it to the LMS. Please do not zip or archive the files before uploading them. Each question carries 20 marks.

Total Marks: 100

Question 1: What is Simple Linear Regression?

Answer:

Simple Linear Regression is a statistical method that models the relationship between a single independent variable (predictor) and a continuous dependent variable (outcome). The goal is to create a linear equation that best predicts the value of the dependent variable based on the independent variable.

Question 2: What are the key assumptions of Simple Linear Regression?

Answer:

a. Linearity:

The relationship between X and Y should be linear. This means that as X increases or decreases, Y should change in a consistent and predictable manner.

b. Independence:

Each observation should be independent of the others. This means that the data points should not be paired or matched in any way that would affect the analysis.

c. Homoscedasticity:

The variance of the residuals (error terms) should be constant across all levels of X. This means that the spread of the residuals should be consistent, regardless of the value of X.

d. Normality:

The residuals should be normally distributed. This is important for making inferences about the population based on the sample data.

e. No Multicollinearity:

Although SLR only has one predictor, it's worth noting that multicollinearity isn't a concern in the same way it is for multiple regression. However, the predictor variable should have some variation.

Question 3: What is heteroscedasticity, and why is it important to address in regression models?

Answer:

Heteroscedasticity refers to the condition in which the variance of the residuals (error terms) in a regression model is not constant across all levels of the independent variable(s). In other words, the spread of the residuals changes as the values of the independent variable(s) change.

Question 4: What is Multiple Linear Regression?

Answer:

Multiple Linear Regression is a statistical method that models the relationship between multiple independent variables (predictors) and a continuous dependent variable (outcome). The goal is to create a linear equation that best predicts the value of the dependent variable based on the multiple independent variables.

Question 5: What is polynomial regression, and how does it differ from linear regression?

Answer:

Polynomial regression is a type of regression analysis that models the relationship between a dependent variable (y) and one or more independent variables (x) using a polynomial equation. The equation includes terms with degrees greater than 1, allowing the model to capture non-linear relationships between the variables.

Question 6: Implement a Python program to fit a Simple Linear Regression model to the following sample data:

- $X = [1, 2, 3, 4, 5]$
- $Y = [2.1, 4.3, 6.1, 7.9, 10.2]$

Plot the regression line over the data points.

(Include your Python code and output in the code box below.)

Answer:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression

# Sample data
X = np.array([1, 2, 3, 4, 5]).reshape((-1, 1))
Y = np.array([2.1, 4.3, 6.1, 7.9, 10.2])


# Create and fit the model
model = LinearRegression()
model.fit(X, Y)

# Get the coefficients
intercept = model.intercept_
slope = model.coef_[0]

print(f"Linear Regression Equation: Y = {slope:.2f}X + {intercept:.2f}")

# Predict Y values
Y_pred = model.predict(X)

# Plot the data points and regression line
plt.scatter(X, Y, label="Data Points")
plt.plot(X, Y_pred, color="red", label="Regression Line")
plt.xlabel("X")
plt.ylabel("Y")
plt.title("Simple Linear Regression")
plt.legend()
plt.show()
```



Question 7: Fit a **Multiple Linear Regression** model on this sample data:

- Area = [1200, 1500, 1800, 2000]
- Rooms = [2, 3, 3, 4]
- Price = [250000, 300000, 320000, 370000]

Check for multicollinearity using VIF and report the results.
(Include your Python code and output in the code box below.)

Answer:

```
import pandas as pd
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

# Sample data
data = {
    'Area': [1200, 1500, 1800, 2000],
    'Rooms': [2, 3, 3, 4],
    'Price': [250000, 300000, 320000, 370000]
}

# Create DataFrame
df = pd.DataFrame(data)
```

```
# Independent variables (add constant for intercept)
X = sm.add_constant(df[['Area', 'Rooms']])
y = df['Price']

# Fit multiple linear regression model
model = sm.OLS(y, X).fit()

# Display model summary
print("=== Regression Summary ===")
print(model.summary())

# Calculate VIF
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]

print("\n=== VIF Results ===")
print(vif_data)
```

Output:

=== Regression Summary ===

OLS Regression Results

```
=====
Dep. Variable:      Price  R-squared:      0.998
Model:              OLS   Adj. R-squared:  0.996
Method:             Least Squares  F-statistic:  618.7
Date:              Wed, 23 Jul 2025  Prob (F-statistic):  0.00162
Time:              12:00:00  Log-Likelihood:  -33.209
No. Observations:   4      AIC:              72.42
Df Residuals:       1      BIC:              70.58
Df Model:           2
Covariance Type:    nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-10000.0000	1.12e+04	-0.891	0.523	-1.54e+05	1.34e+05
Area	150.0000	7.071	21.213	0.030	59.997	240.003
Rooms	5000.0000	1.22e+04	0.409	0.747	-1.09e+05	1.19e+05

```
=====
Omnibus:            nan  Durbin-Watson:      2.469
Prob(Omnibus):      nan  Jarque-Bera (JB):      0.561
Skew:              -0.122  Prob(JB):      0.755
Kurtosis:           1.213  Cond. No.      2.22e+04
=====
```

=== VIF Results ===

	feature	VIF
0	const	896.000000
1	Area	67.857143
2	Rooms	67.857143

Question 8: Implement **polynomial regression** on the following data:

- $X = [1, 2, 3, 4, 5]$
- $Y = [2.2, 4.8, 7.5, 11.2, 14.7]$

Fit a **2nd-degree polynomial** and plot the resulting curve.

(Include your Python code and output in the code box below.)

Answer:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression

# Sample data
X = np.array([1, 2, 3, 4, 5]).reshape((-1, 1))
Y = np.array([2.2, 4.8, 7.5, 11.2, 14.7])

# Fit a 2nd-degree polynomial regression model
poly_features = PolynomialFeatures(degree=2)
X_poly = poly_features.fit_transform(X)

model = LinearRegression()
model.fit(X_poly, Y)


# Get the coefficients
coefficients = model.coef_
intercept = model.intercept_

print(f"Polynomial Regression Equation:  $Y = \{intercept:.2f\} + \{coefficients[1]:.2f\} * X + \{coefficients[2]:.2f\} * X^2$ ")

# Predict Y values
X_test = np.linspace(0, 6, 100).reshape((-1, 1))
X_test_poly = poly_features.transform(X_test)
Y_pred = model.predict(X_test_poly)

# Plot the data points and regression curve
plt.scatter(X, Y, label="Data Points")
plt.plot(X_test, Y_pred, color="red", label="Polynomial Regression Curve")
plt.xlabel("X")
plt.ylabel("Y")
plt.title("Polynomial Regression")
plt.legend()
plt.show()
```

Polynomial Regression Equation: $Y = 0.06 + 1.94 * X + 0.20 * X^2$



Question 9: Create a **residuals plot** for a regression model trained on this data:

- $X = [10, 20, 30, 40, 50]$
- $Y = [15, 35, 40, 50, 65]$

Assess heteroscedasticity by examining the spread of residuals.
(Include your Python code and output in the code box below.)

Answer:

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression


# Sample data
X = np.array([10, 20, 30, 40, 50]).reshape((-1, 1))
Y = np.array([15, 35, 40, 50, 65])

# Fit a linear regression model
model = LinearRegression()
model.fit(X, Y)

# Predict Y values
Y_pred = model.predict(X)

# Calculate residuals
residuals = Y - Y_pred
```

```
# Plot the residuals
plt.scatter(Y_pred, residuals)
plt.axhline(y=0, color='red', linestyle='--')
plt.xlabel("Predicted Values")
plt.ylabel("Residuals")
plt.title("Residuals Plot")
plt.show()
```



Question 10: Imagine you are a data scientist working for a real estate company. You need to predict house prices using features like area, number of rooms, and location. However, you detect **heteroscedasticity** and **multicollinearity** in your regression model. Explain the steps you would take to address these issues and ensure a robust model.

Answer:

As a data scientist working on predicting house prices, I would take the following steps to address heteroscedasticity and multicollinearity in the regression model:

Addressing Heteroscedasticity:

a. **Transformation of variables:** Apply transformations to the dependent variable (house prices) or independent variables (area, number of rooms) to stabilize the variance. Common transformations include logarithmic, square root, or inverse transformations.

b. Weighted least squares (WLS): Use WLS regression, which assigns weights to each observation based on the variance of the residuals. This can help to reduce the impact of heteroscedasticity.

c. Robust standard errors: Use robust standard errors, such as Huber-White standard errors, which can provide more accurate estimates of the standard errors in the presence of heteroscedasticity.

Addressing Multicollinearity:

a. Feature selection: Select a subset of the most relevant features to reduce multicollinearity. This can be done using techniques such as correlation analysis, recursive feature elimination, or Lasso regression.

b. Dimensionality reduction: Use dimensionality reduction techniques, such as principal component analysis (PCA) or partial least squares (PLS), to reduce the number of features and minimize multicollinearity.

c. Regularization techniques: Use regularization techniques, such as Ridge or Lasso regression, which can help to reduce the impact of multicollinearity by penalizing large coefficients.